Artificial Intelligence Nanodegree

Convolutional Neural Networks

Project: Write an Algorithm for a Dog Identification App

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with **'(IMPLEMENTATION)'** in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

> **Note**: Once you have completed all of the code implementations, you need to finalize your work by exporting the iPython Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to \n", "**File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a **'Question X'** header. Carefully read each question and provide thorough answers in the following text boxes that begin with **'Answer:'**. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

> **Note:** Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this IPython notebook.

Why We're Here

In this notebook, you will make the first steps towards developing an algorithm that could be used as part of a mobile or web app. At the end of this project, your code will accept any user-supplied image as input. If a dog is detected in the image, it will provide an estimate of the dog's breed. If a human is detected, it will provide an estimate of the dog breed that is most resembling. The image below displays potential sample output of your finished project (... but we expect that each student's algorithm will behave differently!).

```
hello, dog!
your predicted breed is ...
American Staffordshire terrier
```

In this real-world setting, you will need to piece together a series of models to perform different tasks; for instance, the algorithm that detects humans in an image will be different from the CNN that infers dog breed. There are many points of possible failure, and no perfect algorithm exists. Your imperfect solution will nonetheless create a fun user experience!

The Road Ahead

We break the notebook into separate steps. Feel free to use the links below to navigate the notebook.

- Step 0: Import Datasets
- Step 1: Detect Humans
- Step 2: Detect Dogs
- Step 3: Create a CNN to Classify Dog Breeds (from Scratch)
- Step 4: Use a CNN to Classify Dog Breeds (using Transfer Learning)
- Step 5: Create a CNN to Classify Dog Breeds (using Transfer Learning)
- Step 6: Write your Algorithm
- Step 7: Test Your Algorithm

Step 0: Import Datasets

Import Dog Dataset

In the code cell below, we import a dataset of dog images. We populate a few variables through the use of the `load_files` function from the scikit-learn library:

- `train_files`, `valid_files`, `test_files` - numpy arrays containing file paths to images
- `train_targets`, `valid_targets`, `test_targets` - numpy arrays containing onehot-encoded classification labels

```
In [1]:  from sklearn.datasets import load_files
         from keras.utils import np_utils
         import numpy as np
         from glob import glob

         # define function to load train, test, and validation datasets
         def load_dataset(path):
             data = load_files(path)
             dog_files = np.array(data['filenames'])
             dog_targets = np_utils.to_categorical(np.array(data['target']), 133)
             return dog_files, dog_targets

         # load train, test, and validation datasets
         train_files, train_targets = load_dataset('dogImages/train')
         valid_files, valid_targets = load_dataset('dogImages/valid')
         test_files, test_targets = load_dataset('dogImages/test')

         # load list of dog names
         dog_names = [item[20:-1] for item in sorted(glob("dogImages/train/*/"))]

         # print statistics about the dataset
         print('There are %d total dog categories.' % len(dog_names))
         print('There are %s total dog images.\n' % len(np.hstack([train_files, v
         alid_files, test_files])))
         print('There are %d training dog images.' % len(train_files))
         print('There are %d validation dog images.' % len(valid_files))
         print('There are %d test dog images.'% len(test_files))
```

```
Using TensorFlow backend.

There are 133 total dog categories.
There are 8351 total dog images.

There are 6680 training dog images.
There are 835 validation dog images.
There are 836 test dog images.
```

Import Human Dataset

In the code cell below, we import a dataset of human images, where the file paths are stored in the numpy array
`human_files`.

```
In [2]:  import random
         random.seed(8675309)

         # load filenames in shuffled human dataset
         human_files = np.array(glob("lfw/*/*"))
         random.shuffle(human_files)

         # print statistics about the dataset
         print('There are %d total human images.' % len(human_files))
```

```
There are 13233 total human images.
```

Step 1: Detect Humans

We use OpenCV's implementation of Haar feature-based cascade classifiers (http://docs.opencv.org/trunk/d7/d8b/tutorial_py_face_detection.html) to detect human faces in images. OpenCV provides many pre-trained face detectors, stored as XML files on github (https://github.com/opencv/opencv/tree/master/data/haarcascades). We have downloaded one of these detectors and stored it in the `haarcascades` directory.

In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

In [3]:
```python
import cv2
import matplotlib.pyplot as plt
%matplotlib inline

# extract pre-trained face detector
face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalfa
ce_alt.xml')

# load color (BGR) image
img = cv2.imread(human_files[5])
# convert BGR image to grayscale
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

# find faces in image
faces = face_cascade.detectMultiScale(gray)

# print number of faces detected in the image
print('Number of faces detected:', len(faces))

# get bounding box for each detected face
for (x,y,w,h) in faces:
    # add bounding box to color image
    cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)

# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(cv_rgb)
plt.show()
```
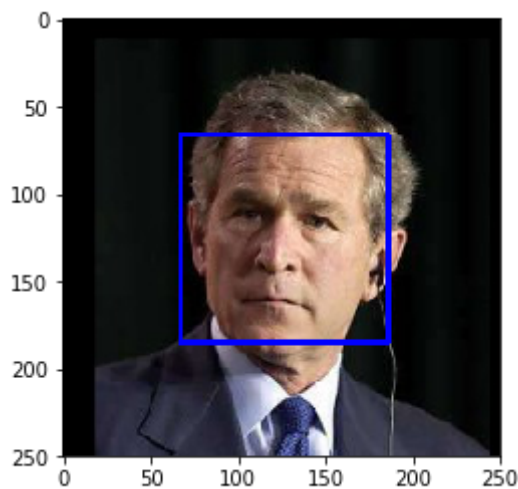
Number of faces detected: 1

```
In [ ]:  def print_img(img_path):
             # extract pre-trained face detector
             face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_front
         alface_alt.xml')

             # load color (BGR) image
             img = cv2.imread(img_path)
             # convert BGR image to grayscale
             gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

             # find faces in image
             faces = face_cascade.detectMultiScale(gray)

             # print number of faces detected in the image
             print('Number of faces detected:', len(faces))

             # get bounding box for each detected face
             for (x,y,w,h) in faces:
                 # add bounding box to color image
                 cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)

             # convert BGR image to RGB for plotting
             cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

             # display the image, along with bounding box
             plt.imshow(cv_rgb)
             plt.show()
             return
```

Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

Write a Human Face Detector

We can use this procedure to write a function that returns `True` if a human face is detected in an image and `False` otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [4]:  # returns "True" if face is detected in image stored at img_path
         def face_detector(img_path):
             img = cv2.imread(img_path)
             gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
             faces = face_cascade.detectMultiScale(gray)
             return len(faces) > 0
```

(IMPLEMENTATION) Assess the Human Face Detector

Question 1: Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

Answer:

Percentage true of human files with human faces 98.0. Percentage true of dog files with human faces 11.0

```
In [5]: human_files_short = human_files[:100]
        dog_files_short = train_files[:100]
        # Do NOT modify the code above this line.

        ## TODO: Test the performance of the face_detector algorithm
        ## on the images in human_files_short and dog_files_short.
        from collections import Counter

        human_count = Counter([face_detector(x) for x in human_files_short])
        print ("Percentage true of human files with human faces " + str(human_co
        unt[True] / (human_count[True]+human_count[False]) * 100))

        dog_human_count = Counter([face_detector(y) for y in dog_files_short])
        print ("Percentage true of dog files with human faces " + str(dog_human_
        count[True] / (dog_human_count[True]+dog_human_count[False]) * 100))
```

```
Percentage true of human files with human faces 98.0
Percentage true of dog files with human faces 11.0
```

```
In [6]: for index,value in enumerate(dog_files_short):
            if face_detector(value) == True:
                print (index)
```

```
0
14
15
21
22
23
24
30
32
63
78
```

Question 2: This algorithmic choice necessitates that we communicate to the user that we accept human images only when they provide a clear view of a face (otherwise, we risk having unneccessarily frustrated users!). In your opinion, is this a reasonable expectation to pose on the user? If not, can you think of a way to detect humans in images that does not necessitate an image with a clearly presented face?

Answer:

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on each of the datasets.

```
In [7]:  ## (Optional) TODO: Report the performance of another
         ## face detection algorithm on the LFW dataset
         ### Feel free to use as many code cells as needed.
```

Step 2: Detect Dogs

In this section, we use a pre-trained ResNet-50 (http://ethereon.github.io/netscope/#/gist/db945b393d40bfa26006) model to detect dogs in images. Our first line of code downloads the ResNet-50 model, along with weights that have been trained on ImageNet (http://www.image-net.org/), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of 1000 categories (https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a). Given an image, this pre-trained ResNet-50 model returns a prediction (derived from the available categories in ImageNet) for the object that is contained in the image.

```
In [15]:  from keras.applications.resnet50 import ResNet50

          # define ResNet50 model
          ResNet50_model = ResNet50(weights='imagenet')
```

Pre-process the Data

When using TensorFlow as backend, Keras CNNs require a 4D array (which we'll also refer to as a 4D tensor) as input, with shape

$$

(\text{nb_samples}, \text{rows}, \text{columns}, \text{channels}), $$

where `nb_samples` corresponds to the total number of images (or samples), and `rows`, `columns`, and `channels` correspond to the number of rows, columns, and channels for each image, respectively.

The `path_to_tensor` function below takes a string-valued file path to a color image as input and returns a 4D tensor suitable for supplying to a Keras CNN. The function first loads the image and resizes it to a square image that is 224×224 pixels. Next, the image is converted to an array, which is then resized to a 4D tensor. In this case, since we are working with color images, each image has three channels. Likewise, since we are processing a single image (or sample), the returned tensor will always have shape

$$(1, 224, 224, 3).$$

The `paths_to_tensor` function takes a numpy array of string-valued image paths as input and returns a 4D tensor with shape

$$(\text{nb\_samples}, 224, 224, 3).$$

Here, `nb_samples` is the number of samples, or number of images, in the supplied array of image paths. It is best to think of `nb_samples` as the number of 3D tensors (where each 3D tensor corresponds to a different image) in your dataset!

```
In [9]:  from keras.preprocessing import image
         from tqdm import tqdm

         def path_to_tensor(img_path):
             # loads RGB image as PIL.Image.Image type
             img = image.load_img(img_path, target_size=(224, 224))
             # convert PIL.Image.Image type to 3D tensor with shape (224, 224, 3)
             x = image.img_to_array(img)
             # convert 3D tensor to 4D tensor with shape (1, 224, 224, 3) and ret
         urn 4D tensor
             return np.expand_dims(x, axis=0)

         def paths_to_tensor(img_paths):
             list_of_tensors = [path_to_tensor(img_path) for img_path in tqdm(img
         _paths)]
             return np.vstack(list_of_tensors)
```

Making Predictions with ResNet-50

Getting the 4D tensor ready for ResNet-50, and for any other pre-trained model in Keras, requires some additional processing. First, the RGB image is converted to BGR by reordering the channels. All pre-trained models have the additional normalization step that the mean pixel (expressed in RGB as $[103.939, 116.779, 123.68]$ and calculated from all pixels in all images in ImageNet) must be subtracted from every pixel in each image. This is implemented in the imported function `preprocess_input`. If you're curious, you can check the code for `preprocess_input` here (https://github.com/fchollet/keras/blob/master/keras/applications/imagenet_utils.py).

Now that we have a way to format our image for supplying to ResNet-50, we are now ready to use the model to extract the predictions. This is accomplished with the `predict` method, which returns an array whose i-th entry is the model's predicted probability that the image belongs to the i-th ImageNet category. This is implemented in the `ResNet50_predict_labels` function below.

By taking the argmax of the predicted probability vector, we obtain an integer corresponding to the model's predicted object class, which we can identify with an object category through the use of this dictionary (https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a).

```
In [10]:  from keras.applications.resnet50 import preprocess_input, decode_predict
          ions

          def ResNet50_predict_labels(img_path):
              # returns prediction vector for image located at img_path
              img = preprocess_input(path_to_tensor(img_path))
              return np.argmax(ResNet50_model.predict(img))
```

Write a Dog Detector

While looking at the dictionary (https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from `'Chihuahua'` to `'Mexican hairless'`. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained ResNet-50 model, we need only check if the `ResNet50_predict_labels` function above returns a value between 151 and 268 (inclusive).

We use these ideas to complete the `dog_detector` function below, which returns `True` if a dog is detected in an image (and `False` if not).

```
In [11]:  ### returns "True" if a dog is detected in the image stored at img_path
          def dog_detector(img_path):
              prediction = ResNet50_predict_labels(img_path)
              return ((prediction <= 268) & (prediction >= 151))
```

(IMPLEMENTATION) Assess the Dog Detector

Question 3: Use the code cell below to test the performance of your `dog_detector` function.

- What percentage of the images in `human_files_short` have a detected dog?
- What percentage of the images in `dog_files_short` have a detected dog?

Answer:

Percentage true of human files with a dog 1.0. Percentage true of dog files with a dog 100.0

```
In [124]:   ### TODO: Test the performance of the dog_detector function
            ### on the images in human_files_short and dog_files_short.

            human_dog_count = Counter([dog_detector(x) for x in human_files_short])
            print ("Percentage true of human files with a dog " + str(human_dog_coun
            t[True] / (human_dog_count[True]+human_dog_count[False]) * 100))

            dog_count = Counter([dog_detector(y) for y in dog_files_short])
            print ("Percentage true of dog files with a dog " + str(dog_count[True]
            / (dog_count[True]+dog_count[False]) * 100))
```

```
            Percentage true of human files with a dog 1.0
            Percentage true of dog files with a dog 100.0
```

```
In [42]:    for index,x in enumerate(human_files_short):
                if dog_detector(x) == True:
                    print (index)
```

```
            71
```

# Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet*!), and you must attain a test accuracy of at least 1%. In Step 5 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

Be careful with adding too many trainable layers! More parameters means longer training, which means you are more likely to need a GPU to accelerate the training process. Thankfully, Keras provides a handy estimate of the time that each epoch is likely to take; you can extrapolate this estimate to figure out how long it will take for your algorithm to train.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have great difficulty in distinguishing between a Brittany and a Welsh Springer Spaniel.

Brittany	Welsh Springer Spaniel
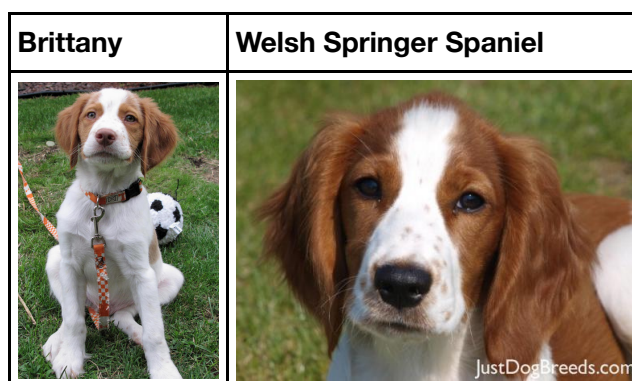
It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

Curly-Coated Retriever	American Water Spaniel

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

Yellow Labrador	Chocolate Labrador	Black Labrador

Yellow Labrador	Chocolate Labrador	Black Labrador

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imabalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

Pre-process the Data

```
In [17]:   from PIL import ImageFile
           ImageFile.LOAD_TRUNCATED_IMAGES = True

           # pre-process the data for Keras
           train_tensors = paths_to_tensor(train_files).astype('float32')/255
           valid_tensors = paths_to_tensor(valid_files).astype('float32')/255
           test_tensors = paths_to_tensor(test_files).astype('float32')/255
```

```
100%|██████████| 6680/6680 [01:10<00:00, 94.28it/s]
100%|██████████| 835/835 [00:08<00:00, 95.51it/s]
100%|██████████| 836/836 [00:08<00:00, 95.41it/s]
```

(IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. At the end of your code cell block, summarize the layers of your model by executing the line:

```
model.summary()
```

We have imported some Python modules to get you started, but feel free to import as many modules as you need. If you end up getting stuck, here's a hint that specifies a model that trains relatively fast on CPU and attains >1% test accuracy in 5 epochs:

```
Layer (type)                 Output Shape              Param #
=================================================================
conv2d_1 (Conv2D)            (None, 223, 223, 16)      208
_____
max_pooling2d_1 (MaxPooling2 (None, 111, 111, 16)      0
_____
conv2d_2 (Conv2D)            (None, 110, 110, 32)      2080
_____
max_pooling2d_2 (MaxPooling2 (None, 55, 55, 32)        0
_____
conv2d_3 (Conv2D)            (None, 54, 54, 64)        8256
_____
max_pooling2d_3 (MaxPooling2 (None, 27, 27, 64)        0
_____
global_average_pooling2d_1 ( (None, 64)                0
_____
dense_1 (Dense)              (None, 133)               8645
=================================================================
Total params: 19,189.0
Trainable params: 19,189.0
Non-trainable params: 0.0
_____
```

INPUT

CONV

POOL

CONV

POOL

CONV

POOL

GAP

DENSE

Question 4: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. If you chose to use the hinted architecture above, describe why you think that CNN architecture should work well for the image classification task.

Answer: I chose to use the CNN architecture specified because it seemed like an appropriate amount of convolutional and pooling layers which ultimately led to the final dense output layer. From my research, 3 layers of convultions (with relu activations) and max pooling seemed to be a standard starting architecture for a CNN.

In [43]:
```python
print('x_train shape:', train_tensors.shape)
train_tensors_sub = train_tensors[0:300]
print('x_train shape:', train_tensors_sub.shape)
```

```
x_train shape: (6680, 224, 224, 3)
x_train shape: (300, 224, 224, 3)
```

In [63]:
```python
# from keras.layers import Conv2D, MaxPooling2D, GlobalAveragePooling2D
from keras.layers import Dropout, Flatten, Dense
from keras.models import Sequential

model = Sequential()

### TODO: Define your architecture.

model.add(Conv2D(filters = 16, kernel_size = 2, strides = 2,
        padding = 'same', activation = 'relu',
        input_shape = (224,224,3)))
model.add(MaxPooling2D(pool_size=2))
model.add(Conv2D(filters = 32, kernel_size = 2, strides = 1,
        padding = 'VALID', activation = 'relu'))
model.add(MaxPooling2D(pool_size=2))

model.add(Conv2D(filters = 64, kernel_size = 2, strides = 1,
        padding = 'VALID', activation = 'relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(GlobalAveragePooling2D())
model.add(Dense(133,activation='relu'))


model.summary()
```

```
_____
Layer (type)                 Output Shape              Param #
=================================================================
conv2d_77 (Conv2D)           (None, 112, 112, 16)      208
_____
max_pooling2d_78 (MaxPooling (None, 56, 56, 16)        0
_____
conv2d_78 (Conv2D)           (None, 55, 55, 32)        2080
_____
max_pooling2d_79 (MaxPooling (None, 27, 27, 32)        0
_____
conv2d_79 (Conv2D)           (None, 26, 26, 64)        8256
_____
max_pooling2d_80 (MaxPooling (None, 13, 13, 64)        0
_____
global_average_pooling2d_25  (None, 64)                0
_____
dense_3 (Dense)              (None, 133)               8645
=================================================================
Total params: 19,189
Trainable params: 19,189
Non-trainable params: 0
_____
```

Compile the Model

```
In [64]: model.compile(optimizer='rmsprop', loss='categorical_crossentropy', metr
ics=['accuracy'])
```

(IMPLEMENTATION) Train the Model

Train your model in the code cell below. Use model checkpointing to save the model that attains the best validation loss.

You are welcome to augment the training data (https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html), but this is not a requirement.

```
In [65]:   from keras.callbacks import ModelCheckpoint

           ### TODO: specify the number of epochs that you would like to use to tra
           in the model.

           epochs = 5

           ### Do NOT modify the code below this line.

           checkpointer = ModelCheckpoint(filepath='saved_models/weights.best.from_
           scratch.hdf5',
                                          verbose=1, save_best_only=True)

           model.fit(train_tensors, train_targets,
                     validation_data=(valid_tensors, valid_targets),
                     epochs=epochs, batch_size=20, callbacks=[checkpointer], verbos
           e=1)
```

```
Train on 6680 samples, validate on 835 samples
Epoch 1/5
6660/6680 [=============================>.] - ETA: 0s - loss: 12.1693 -
acc: 0.0078Epoch 00001: val_loss improved from inf to 11.97162, saving
model to saved_models/weights.best.from_scratch.hdf5
6680/6680 [==============================] - 72s 11ms/step - loss: 12.1
593 - acc: 0.0078 - val_loss: 11.9716 - val_acc: 0.0096
Epoch 2/5
6660/6680 [=============================>.] - ETA: 0s - loss: 12.4451 -
acc: 0.0084Epoch 00002: val_loss did not improve
6680/6680 [==============================] - 60s 9ms/step - loss: 12.44
28 - acc: 0.0084 - val_loss: 12.9738 - val_acc: 0.0096
Epoch 3/5
6660/6680 [=============================>.] - ETA: 0s - loss: 12.6795 -
acc: 0.0093Epoch 00003: val_loss did not improve
6680/6680 [==============================] - 67s 10ms/step - loss: 12.6
765 - acc: 0.0094 - val_loss: 12.5376 - val_acc: 0.0072
Epoch 4/5
6660/6680 [=============================>.] - ETA: 0s - loss: 12.5590 -
acc: 0.0075Epoch 00004: val_loss did not improve
6680/6680 [==============================] - 71s 11ms/step - loss: 12.5
641 - acc: 0.0075 - val_loss: 12.3947 - val_acc: 0.0072
Epoch 5/5
6660/6680 [=============================>.] - ETA: 0s - loss: 13.7554 -
acc: 0.0101Epoch 00005: val_loss did not improve
6680/6680 [==============================] - 74s 11ms/step - loss: 13.7
624 - acc: 0.0100 - val_loss: 15.2318 - val_acc: 0.0108
```

Out[65]:   <keras.callbacks.History at 0x1a2ed39ac8>

Load the Model with the Best Validation Loss

```
In [66]:   model.load_weights('saved_models/weights.best.from_scratch.hdf5')
```

Test the Model

Try out your model on the test dataset of dog images. Ensure that your test accuracy is greater than 1%.

```
In [67]:  # get index of predicted dog breed for each image in test set
          dog_breed_predictions = [np.argmax(model.predict(np.expand_dims(tensor,
          axis=0))) for tensor in test_tensors]

          # report test accuracy
          test_accuracy = 100*np.sum(np.array(dog_breed_predictions)==np.argmax(te
          st_targets, axis=1))/len(dog_breed_predictions)
          print('Test accuracy: %.4f%%' % test_accuracy)
```

```
Test accuracy: 1.0766%
```

Step 4: Use a CNN to Classify Dog Breeds

To reduce training time without sacrificing accuracy, we show you how to train a CNN using transfer learning. In the following step, you will get a chance to use transfer learning to train your own CNN.

Obtain Bottleneck Features

```
In [69]:  bottleneck_features = np.load('bottleneck_features/DogVGG16Data.npz')
          train_VGG16 = bottleneck_features['train']
          valid_VGG16 = bottleneck_features['valid']
          test_VGG16 = bottleneck_features['test']
```

Model Architecture

The model uses the the pre-trained VGG-16 model as a fixed feature extractor, where the last convolutional output of VGG-16 is fed as input to our model. We only add a global average pooling layer and a fully connected layer, where the latter contains one node for each dog category and is equipped with a softmax.

In [70]:
```
VGG16_model = Sequential()
VGG16_model.add(GlobalAveragePooling2D(input_shape=train_VGG16.shape[1
:]))
VGG16_model.add(Dense(133, activation='softmax'))

VGG16_model.summary()
```

```
_____
Layer (type)                 Output Shape              Param #
=================================================================
global_average_pooling2d_26  (None, 512)               0
_____
dense_4 (Dense)              (None, 133)               68229
=================================================================
Total params: 68,229
Trainable params: 68,229
Non-trainable params: 0
_____
```

Compile the Model

In [71]:
```
VGG16_model.compile(loss='categorical_crossentropy', optimizer='rmsprop'
, metrics=['accuracy'])
```

Train the Model

```
In [72]: checkpointer = ModelCheckpoint(filepath='saved_models/weights.best.VGG1
         6.hdf5',
                                        verbose=1, save_best_only=True)

         VGG16_model.fit(train_VGG16, train_targets,
                 validation_data=(valid_VGG16, valid_targets),
                 epochs=20, batch_size=20, callbacks=[checkpointer], verbose=1)
```

```
Train on 6680 samples, validate on 835 samples
Epoch 1/20
6600/6680 [=============================>.] - ETA: 0s - loss: 12.1121 -
acc: 0.1347Epoch 00001: val_loss improved from inf to 10.76260, saving
model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 4s 671us/step - loss: 12.0
832 - acc: 0.1361 - val_loss: 10.7626 - val_acc: 0.1964
Epoch 2/20
6480/6680 [=============================>.] - ETA: 0s - loss: 9.9756 - a
cc: 0.2971Epoch 00002: val_loss improved from 10.76260 to 10.06791, sav
ing model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 2s 257us/step - loss: 9.97
41 - acc: 0.2981 - val_loss: 10.0679 - val_acc: 0.2814
Epoch 3/20
6580/6680 [=============================>.] - ETA: 0s - loss: 9.4198 - a
cc: 0.3591Epoch 00003: val_loss improved from 10.06791 to 9.67521, savi
ng model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 2s 241us/step - loss: 9.41
47 - acc: 0.3596 - val_loss: 9.6752 - val_acc: 0.3269
Epoch 4/20
6620/6680 [=============================>.] - ETA: 0s - loss: 9.1651 - a
cc: 0.3956Epoch 00004: val_loss improved from 9.67521 to 9.56776, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 2s 244us/step - loss: 9.17
40 - acc: 0.3951 - val_loss: 9.5678 - val_acc: 0.3257
Epoch 5/20
6540/6680 [=============================>.] - ETA: 0s - loss: 9.0625 - a
cc: 0.4109Epoch 00005: val_loss did not improve
6680/6680 [==============================] - 2s 250us/step - loss: 9.06
06 - acc: 0.4109 - val_loss: 9.5846 - val_acc: 0.3317
Epoch 6/20
6560/6680 [=============================>.] - ETA: 0s - loss: 8.9652 - a
cc: 0.4220Epoch 00006: val_loss improved from 9.56776 to 9.39897, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 2s 247us/step - loss: 8.97
34 - acc: 0.4216 - val_loss: 9.3990 - val_acc: 0.3521
Epoch 7/20
6460/6680 [=============================>.] - ETA: 0s - loss: 8.8520 - a
cc: 0.4330Epoch 00007: val_loss did not improve
6680/6680 [==============================] - 2s 262us/step - loss: 8.84
23 - acc: 0.4332 - val_loss: 9.4781 - val_acc: 0.3485
Epoch 8/20
6540/6680 [=============================>.] - ETA: 0s - loss: 8.6862 - a
cc: 0.4425Epoch 00008: val_loss improved from 9.39897 to 9.30088, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 2s 246us/step - loss: 8.67
47 - acc: 0.4428 - val_loss: 9.3009 - val_acc: 0.3521
Epoch 9/20
6620/6680 [=============================>.] - ETA: 0s - loss: 8.6157 - a
cc: 0.4527Epoch 00009: val_loss improved from 9.30088 to 9.13408, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 2s 245us/step - loss: 8.60
59 - acc: 0.4534 - val_loss: 9.1341 - val_acc: 0.3689
Epoch 10/20
6540/6680 [=============================>.] - ETA: 0s - loss: 8.4752 - a
cc: 0.4624Epoch 00010: val_loss improved from 9.13408 to 9.06201, savin
g model to saved_models/weights.best.VGG16.hdf5
```

```
6680/6680 [==============================] - 2s 250us/step - loss: 8.49
70 - acc: 0.4611 - val_loss: 9.0620 - val_acc: 0.3737
Epoch 11/20
6600/6680 [=============================>.] - ETA: 0s - loss: 8.4356 - a
cc: 0.4677Epoch 00011: val_loss improved from 9.06201 to 9.00281, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 2s 256us/step - loss: 8.42
39 - acc: 0.4686 - val_loss: 9.0028 - val_acc: 0.3832
Epoch 12/20
6660/6680 [=============================>.] - ETA: 0s - loss: 8.4096 - a
cc: 0.4728Epoch 00012: val_loss improved from 9.00281 to 8.93574, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 2s 251us/step - loss: 8.40
86 - acc: 0.4729 - val_loss: 8.9357 - val_acc: 0.3952
Epoch 13/20
6480/6680 [=============================>.] - ETA: 0s - loss: 8.3833 - a
cc: 0.4733Epoch 00013: val_loss improved from 8.93574 to 8.89271, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 2s 228us/step - loss: 8.37
83 - acc: 0.4735 - val_loss: 8.8927 - val_acc: 0.3725
Epoch 14/20
6580/6680 [=============================>.] - ETA: 0s - loss: 8.0745 - a
cc: 0.4787Epoch 00014: val_loss improved from 8.89271 to 8.60975, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 2s 258us/step - loss: 8.05
89 - acc: 0.4793 - val_loss: 8.6097 - val_acc: 0.3880
Epoch 15/20
6540/6680 [=============================>.] - ETA: 0s - loss: 7.7546 - a
cc: 0.5003Epoch 00015: val_loss improved from 8.60975 to 8.53105, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 2s 286us/step - loss: 7.78
05 - acc: 0.4988 - val_loss: 8.5310 - val_acc: 0.4108
Epoch 16/20
6500/6680 [=============================>.] - ETA: 0s - loss: 7.7155 - a
cc: 0.5097Epoch 00016: val_loss improved from 8.53105 to 8.41098, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 2s 312us/step - loss: 7.71
68 - acc: 0.5093 - val_loss: 8.4110 - val_acc: 0.4096
Epoch 17/20
6560/6680 [=============================>.] - ETA: 0s - loss: 7.6856 - a
cc: 0.5140Epoch 00017: val_loss did not improve
6680/6680 [==============================] - 2s 275us/step - loss: 7.68
96 - acc: 0.5138 - val_loss: 8.4381 - val_acc: 0.4036
Epoch 18/20
6620/6680 [=============================>.] - ETA: 0s - loss: 7.5907 - a
cc: 0.5183Epoch 00018: val_loss improved from 8.41098 to 8.31979, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 2s 330us/step - loss: 7.60
63 - acc: 0.5174 - val_loss: 8.3198 - val_acc: 0.4180
Epoch 19/20
6480/6680 [=============================>.] - ETA: 0s - loss: 7.4998 - a
cc: 0.5252Epoch 00019: val_loss improved from 8.31979 to 8.31052, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 2s 277us/step - loss: 7.51
19 - acc: 0.5246 - val_loss: 8.3105 - val_acc: 0.4108
Epoch 20/20
6440/6680 [=============================>..] - ETA: 0s - loss: 7.4695 - a
```

```
cc: 0.5300Epoch 00020: val_loss improved from 8.31052 to 8.28084, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 2s 251us/step - loss: 7.48
03 - acc: 0.5290 - val_loss: 8.2808 - val_acc: 0.4180
```

Out[72]: <keras.callbacks.History at 0x1a295cfeb8>

Load the Model with the Best Validation Loss

In [73]:
```python
VGG16_model.load_weights('saved_models/weights.best.VGG16.hdf5')
```

Test the Model

Now, we can use the CNN to test how well it identifies breed within our test dataset of dog images. We print the test accuracy below.

In [74]:
```python
# get index of predicted dog breed for each image in test set
VGG16_predictions = [np.argmax(VGG16_model.predict(np.expand_dims(featur
e, axis=0))) for feature in test_VGG16]

# report test accuracy
test_accuracy = 100*np.sum(np.array(VGG16_predictions)==np.argmax(test_t
argets, axis=1))/len(VGG16_predictions)
print('Test accuracy: %.4f%%' % test_accuracy)
```

```
Test accuracy: 43.5407%
```

Predict Dog Breed with the Model

In [77]:
```python
from extract_bottleneck_features import *

def VGG16_predict_breed(img_path):
    # extract bottleneck features
    bottleneck_feature = extract_VGG16(path_to_tensor(img_path))
    # obtain predicted vector
    predicted_vector = VGG16_model.predict(bottleneck_feature)
    # return dog breed that is predicted by the model
    return dog_names[np.argmax(predicted_vector)]
```

Step 5: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

In Step 4, we used transfer learning to create a CNN using VGG-16 bottleneck features. In this section, you must use the bottleneck features from a different pre-trained model. To make things easier for you, we have pre-computed the features for all of the networks that are currently available in Keras:

- VGG-19 (https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogVGG19Data.npz) bottleneck features
- ResNet-50 (https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogResnet50Data.npz) bottleneck features
- Inception (https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogInceptionV3Data.npz) bottleneck features
- Xception (https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogXceptionData.npz) bottleneck features

The files are encoded as such:

```
Dog{network}Data.npz
```

where `{network}`, in the above filename, can be one of `VGG19`, `Resnet50`, `InceptionV3`, or `Xception`. Pick one of the above architectures, download the corresponding bottleneck features, and store the downloaded file in the `bottleneck_features/` folder in the repository.

(IMPLEMENTATION) Obtain Bottleneck Features

In the code block below, extract the bottleneck features corresponding to the train, test, and validation sets by running the following:

```
bottleneck_features = np.load('bottleneck_features/Dog{network}Data.npz')
train_{network} = bottleneck_features['train']
valid_{network} = bottleneck_features['valid']
test_{network} = bottleneck_features['test']
```

```
In [78]:   ### TODO: Obtain bottleneck features from another pre-trained CNN.
           bottleneck_features = np.load('bottleneck_features/DogResnet50Data.npz')
           train_resnet50 = bottleneck_features['train']
           valid_resnet50 = bottleneck_features['valid']
           test_resnet50 = bottleneck_features['test']
```

(IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. At the end of your code cell block, summarize the layers of your model by executing the line:

```
<your model's name>.summary()
```

Question 5: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

Answer: My first step was to add a layer between the resnet50 model and then add the final layer with a softmax.

```
In [84]:  ### TODO: Define your architecture.
          resnet50_model = Sequential()
          resnet50_model.add(GlobalAveragePooling2D(input_shape=train_resnet50.sha
          pe[1:]))
          resnet50_model.add(Dense(133, activation='softmax'))

          resnet50_model.summary()
```

```
_____
Layer (type)                     Output Shape              Param #
=================================================================
global_average_pooling2d_29  (None, 2048)                 0
_____
dense_7 (Dense)                  (None, 133)               272517
=================================================================
Total params: 272,517
Trainable params: 272,517
Non-trainable params: 0
_____
```

(IMPLEMENTATION) Compile the Model

```
In [85]:  ### TODO: Compile the model.
          resnet50_model.compile(loss='categorical_crossentropy', optimizer='rmspr
          op', metrics=['accuracy'])
```

(IMPLEMENTATION) Train the Model

Train your model in the code cell below. Use model checkpointing to save the model that attains the best validation loss.

You are welcome to augment the training data (https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html), but this is not a requirement.

In [88]:
```python
### TODO: Train the model.
checkpointer = ModelCheckpoint(filepath='saved_models/weights.best.resne
t50.hdf5',
                                verbose=1, save_best_only=True)

resnet50_model.fit(train_resnet50, train_targets,
        validation_data=(valid_resnet50, valid_targets),
        epochs=20, batch_size=20, callbacks=[checkpointer], verbose=1)
```

```
Train on 6680 samples, validate on 835 samples
Epoch 1/20
6600/6680 [=============================>.] - ETA: 0s - loss: 0.0998 - a
cc: 0.9680Epoch 00001: val_loss improved from inf to 0.62786, saving mo
del to saved_models/weights.best.resnet50.hdf5
6680/6680 [==============================] - 2s 373us/step - loss: 0.09
98 - acc: 0.9680 - val_loss: 0.6279 - val_acc: 0.8168
Epoch 2/20
6500/6680 [=============================>.] - ETA: 0s - loss: 0.0673 - a
cc: 0.9803Epoch 00002: val_loss did not improve
6680/6680 [==============================] - 2s 270us/step - loss: 0.06
71 - acc: 0.9801 - val_loss: 0.6905 - val_acc: 0.8168
Epoch 3/20
6540/6680 [=============================>.] - ETA: 0s - loss: 0.0514 - a
cc: 0.9850Epoch 00003: val_loss did not improve
6680/6680 [==============================] - 2s 315us/step - loss: 0.05
09 - acc: 0.9853 - val_loss: 0.7749 - val_acc: 0.8036
Epoch 4/20
6560/6680 [=============================>.] - ETA: 0s - loss: 0.0373 - a
cc: 0.9883Epoch 00004: val_loss did not improve
6680/6680 [==============================] - 2s 320us/step - loss: 0.03
70 - acc: 0.9883 - val_loss: 0.7345 - val_acc: 0.8180
Epoch 5/20
6560/6680 [=============================>.] - ETA: 0s - loss: 0.0285 - a
cc: 0.9922Epoch 00005: val_loss did not improve
6680/6680 [==============================] - 2s 322us/step - loss: 0.02
87 - acc: 0.9922 - val_loss: 0.7483 - val_acc: 0.8240
Epoch 6/20
6620/6680 [=============================>.] - ETA: 0s - loss: 0.0229 - a
cc: 0.9938Epoch 00006: val_loss did not improve
6680/6680 [==============================] - 2s 328us/step - loss: 0.02
44 - acc: 0.9937 - val_loss: 0.7658 - val_acc: 0.8228
Epoch 7/20
6560/6680 [=============================>.] - ETA: 0s - loss: 0.0184 - a
cc: 0.9948Epoch 00007: val_loss did not improve
6680/6680 [==============================] - 3s 456us/step - loss: 0.01
83 - acc: 0.9948 - val_loss: 0.7879 - val_acc: 0.8323
Epoch 8/20
6660/6680 [=============================>.] - ETA: 0s - loss: 0.0141 - a
cc: 0.9962Epoch 00008: val_loss did not improve
6680/6680 [==============================] - 3s 445us/step - loss: 0.01
41 - acc: 0.9963 - val_loss: 0.8013 - val_acc: 0.8287
Epoch 9/20
6620/6680 [=============================>.] - ETA: 0s - loss: 0.0133 - a
cc: 0.9973Epoch 00009: val_loss did not improve
6680/6680 [==============================] - 2s 348us/step - loss: 0.01
32 - acc: 0.9973 - val_loss: 0.8077 - val_acc: 0.8263
Epoch 10/20
6660/6680 [=============================>.] - ETA: 0s - loss: 0.0118 - a
cc: 0.9974Epoch 00010: val_loss did not improve
6680/6680 [==============================] - 2s 371us/step - loss: 0.01
17 - acc: 0.9975 - val_loss: 0.7894 - val_acc: 0.8323
Epoch 11/20
6640/6680 [=============================>.] - ETA: 0s - loss: 0.0083 - a
cc: 0.9976Epoch 00011: val_loss did not simprove
6680/6680 [==============================] - 3s 424us/step - loss: 0.00
83 - acc: 0.9976 - val_loss: 0.8918 - val_acc: 0.8311
```

```
Epoch 12/20
6620/6680 [============================>.] - ETA: 0s - loss: 0.0072 - a
cc: 0.9986Epoch 00012: val_loss did not improve
6680/6680 [==============================] - 3s 456us/step - loss: 0.00
72 - acc: 0.9987 - val_loss: 0.8853 - val_acc: 0.8144
Epoch 13/20
6580/6680 [============================>.] - ETA: 0s - loss: 0.0069 - a
cc: 0.9982Epoch 00013: val_loss did not improve
6680/6680 [==============================] - 3s 438us/step - loss: 0.00
68 - acc: 0.9982 - val_loss: 0.8799 - val_acc: 0.8263
Epoch 14/20
6540/6680 [============================>.] - ETA: 0s - loss: 0.0059 - a
cc: 0.9986Epoch 00014: val_loss did not improve
6680/6680 [==============================] - 2s 328us/step - loss: 0.00
58 - acc: 0.9987 - val_loss: 0.9106 - val_acc: 0.8228
Epoch 15/20
6520/6680 [============================>.] - ETA: 0s - loss: 0.0062 - a
cc: 0.9983Epoch 00015: val_loss did not improve
6680/6680 [==============================] - 3s 379us/step - loss: 0.00
61 - acc: 0.9984 - val_loss: 0.9610 - val_acc: 0.8204
Epoch 16/20
6540/6680 [============================>.] - ETA: 0s - loss: 0.0047 - a
cc: 0.9985Epoch 00016: val_loss did not improve
6680/6680 [==============================] - 2s 320us/step - loss: 0.00
58 - acc: 0.9984 - val_loss: 0.9988 - val_acc: 0.8299
Epoch 17/20
6640/6680 [============================>.] - ETA: 0s - loss: 0.0063 - a
cc: 0.9983Epoch 00017: val_loss did not improve
6680/6680 [==============================] - 2s 331us/step - loss: 0.00
63 - acc: 0.9984 - val_loss: 0.9911 - val_acc: 0.8263
Epoch 18/20
6660/6680 [============================>.] - ETA: 0s - loss: 0.0044 - a
cc: 0.9989Epoch 00018: val_loss did not improve
6680/6680 [==============================] - 3s 381us/step - loss: 0.00
44 - acc: 0.9990 - val_loss: 0.9761 - val_acc: 0.8299
Epoch 19/20
6560/6680 [============================>.] - ETA: 0s - loss: 0.0049 - a
cc: 0.9991Epoch 00019: val_loss did not improve
6680/6680 [==============================] - 2s 369us/step - loss: 0.00
49 - acc: 0.9991 - val_loss: 1.0422 - val_acc: 0.8216
Epoch 20/20
6660/6680 [============================>.] - ETA: 0s - loss: 0.0055 - a
cc: 0.9989Epoch 00020: val_loss did not improve
6680/6680 [==============================] - 3s 469us/step - loss: 0.00
55 - acc: 0.9990 - val_loss: 1.0222 - val_acc: 0.8263
```

Out[88]: <keras.callbacks.History at 0x1a2f72c048>

(IMPLEMENTATION) Load the Model with the Best Validation Loss

```
In [89]: ### TODO: Load the model weights with the best validation loss.
         resnet50_model.load_weights('saved_models/weights.best.resnet50.hdf5')
```

(IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Ensure that your test accuracy is greater than 60%.

```
In [91]:  ### TODO: Calculate classification accuracy on the test dataset.
          # get index of predicted dog breed for each image in test set
          resnet50_predictions = [np.argmax(resnet50_model.predict
                                        (np.expand_dims(feature, axis=0)))
                         for feature in test_resnet50]

          # report test accuracy
          test_accuracy = 100*np.sum(np.array(resnet50_predictions)==
                               np.argmax(test_targets, axis=1))/len(resnet50
          _predictions)
          print('Test accuracy: %.4f%%' % test_accuracy)
```

```
Test accuracy: 81.8182%
```

(IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (`Affenpinscher`, `Afghan_hound`, etc) that is predicted by your model.

Similar to the analogous function in Step 5, your function should have three steps:

1. Extract the bottleneck features corresponding to the chosen CNN model.
2. Supply the bottleneck features as input to the model to return the predicted vector. Note that the argmax of this prediction vector gives the index of the predicted dog breed.
3. Use the `dog_names` array defined in Step 0 of this notebook to return the corresponding breed.

The functions to extract the bottleneck features can be found in `extract_bottleneck_features.py`, and they have been imported in an earlier code cell. To obtain the bottleneck features corresponding to your chosen CNN architecture, you need to use the function

```
extract_{network}
```

where `{network}`, in the above filename, should be one of `VGG19`, `Resnet50`, `InceptionV3`, or `Xception`.

```
In [92]:  #test cell
```

```
Out[92]:  array([[[  3.84162404e-02,   5.12999177e-01,   2.01398935e-02, ...,
                    5.44046640e-01,   5.10268088e-04,   1.25810936e-01]]], dtype
          =float32)
```

```
In [104]: ### TODO: Write a function that takes a path to an image as input
          ### and returns the dog breed that is predicted by the model.

          from extract_bottleneck_features import *

          def resnet50_predict_breed(img_path):
              # extract bottleneck features
              bottleneck_feature = extract_Resnet50(path_to_tensor(img_path))
              # obtain predicted vector
              predicted_vector = resnet50_model.predict(bottleneck_feature)
              # return dog breed that is predicted by the model
              return dog_names[np.argmax(predicted_vector)]
```

Step 6: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then,

- if a **dog** is detected in the image, return the predicted breed.
- if a **human** is detected in the image, return the resembling dog breed.
- if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `dog_detector` functions developed above. You are **required** to use your CNN from Step 5 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!

```
hello, human!
```



```
You look like a ...
Chinese_shar-pei
```

(IMPLEMENTATION) Write your Algorithm

In [105]:
```python
### TODO: Write your algorithm.
### Feel free to use as many code cells as needed.
def dog_or_human(img_path):
    if face_detector(img_path) == True:
        print ("Hello darling")
        print_img(img_path)
        return "You look like a .... " + resnet50_predict_breed(img_path
)
    elif dog_detector(img_path) == True:
        return "This dog is a ..... " + resnet50_predict_breed(img_path)
    else:
        return "Please input a valid picture of a dog or human"
```

Step 7: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that **you** look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

(IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

Question 6: Is the output better than you expected :) ? Or worse :(? Provide at least three possible points of improvement for your algorithm.

Answer: My results are better than I expected. Leveraging the resnet50 features have really boosted the accuracy and the practice of uncovering this via transfer learning was incredible.

I think I could improve my model by

1) Image augmentation so that the model will adapt to different angles/rotations of the same dog and will classify correctly. Plus more training data could help improve results.

2) Try different activation functions such as PReLU or RReLU as they have been identified to outperform ReLU.

3) Find the best combination and order of convolution and pooling layers before submitting the predicted class. Also experimenting with different methods such as batch normalization and adding it into our architecture.

In [106]:
```
## TODO: Execute your algorithm from Step 6 on
## at least 6 images on your computer.
## Feel free to use as many code cells as needed.
dog_or_human('images/frenchbulldog.jpg')
```

Downloading data from https://github.com/fchollet/deep-learning-models/
releases/download/v0.2/resnet50_weights_tf_dim_ordering_tf_kernels_noto
p.h5
94658560/94653016 [==============================] - 55s 1us/step

Out[106]: 'This dog is a ..... French_bulldog'

In [120]:
```
dog_or_human('images/headshot_2.jpg')
```

Hello darling
Number of faces detected: 2

Out[120]: 'You look like a .... Silky_terrier'

In [108]:
```
dog_or_human('images/grumpycat.jpeg')
```

Out[108]: 'Please input a valid picture of a dog or human'

In [114]: `dog_or_human('images/donald.jpg')`

```
Hello darling
Number of faces detected: 1
```


Out[114]: `'You look like a .... Bearded_collie'`

In [110]: `dog_or_human('images/Labrador_retriever_06449.jpg')`

Out[110]: `'This dog is a ..... Labrador_retriever'`

In [111]: `dog_or_human('images/Brittany_02625.jpg')`

Out[111]: `'This dog is a ..... Brittany'`